WORKSTREAM PROPOSAL A3 · SAY-DO LANE · 2026-11-05

Say–do workstreams: what to run next, what each would license, and what it costs

Built from `say-do/reports/A3_say-do_workstreams_roi.md` by `analysis/build_report_pdf.py`. Every judged number cited here is a property of the instrument under named, unvalidated language-model judges; models are named in sealed census order and never ranked.

Contents

Summary
1. How the workstreams are scored
2. Tranche 1: existing data, $0
   - W4. Disclosed versus undisclosed compromise, and record-call order
   - W15. The middle-plan sensitivity table and a cluster bootstrap
   - W7a. Charter form and placement
   - W16. The related-work section, the regime table and the ceiling caveat
   - W17. The judge-ceiling table, and AB6 extended
   - W12. The extraction-artefact check
3. Tranche 1: small metered arms on the Schwartz agentic items
   - W1. The routing arm
   - W2. The priority-reversed and factorial twins
   - W3. The option-set manipulation
   - W8. Advance self-prediction of the middle plan
   - W5. The expert read, extended
4. Tranche 2: arms that extend the finding
   - W11. Post-capitulation re-elicitation on the sustained rung
   - W6. Evaluation awareness as a covariate
   - W9. A reasoning-effort arm
   - W14. The planned owner-machine items
5. Tranche 3: larger arms
   - W10. Cross-benchmark concordance
   - W18. The release package
   - W13. A human comparison arm
   - W7b. Charter placement arm (conditional)
6. Sequencing
7. What the tranches would let the paper say

Workstream proposal (A3) · 2026-11-05 · Multi-Alignment / Moral_Psychology-Paper · source: the open items of findings paper v22 (§S25), `TASKS.md`, `PAPER_PLAN.md` §5, `plans/GAP_CLOSURE_PLAN_20261028.md`, `plans/SPEND_LEDGER.md`, the adversarial council review of 2026-10-31, the say–do corpus's follow-up tables, and the cross-readings of A1 · costs are anchored to realised figures and every paid step still needs its own grant under lane law · PROVISIONAL · judges UNVALIDATED · no ranking.

Summary

Eighteen workstreams are proposed for the say–do line, scored on one rubric and ranked by return on investment. Six run here at $0 on tracked files or existing run trees and should start first; they include the disclosed-versus-undisclosed split of the say-not-do episodes, the middle-plan sensitivity table, the charter-form tabulation, the cluster bootstrap, the related-work section and the judge-ceiling table. Four are small metered arms on the sixteen Schwartz agentic items that would turn the paper's positioning into mechanism for under $150 in all: the routing arm (does the action move when it must pass through the record), the priority-reversed and factorial twins, the option-set manipulation and the advance self-prediction probe. One, the expert read of the reasons channel, is the only item that converts "argue around the judges" into evidence and needs two readers and about ten hours of their time. The remainder are larger (a cross-benchmark concordance arm, a human comparison arm, a reasoning-effort arm) or belong to streams already planned (the frontier repeats on the command-line routes, the deferred gpt-5.6-sol re-judge). The recommended first tranche is

the six $0 items plus W1, W2, W3, W8 and W5, at about $75 metered (cap $140) and seven person-weeks, which would let a v23 of the findings paper, or the say–do paper of A2, make a mechanism claim and an option-set claim it cannot make today.

Table 1. The eighteen workstreams at a glance. Impact is 1–5 on the rubric of §1 and is shown with the risk; cost units are person-days plus metered dollars divided by 25; ROI is impact × 10 ÷ cost units; tranche 1 is the recommended first set. Every metered figure is an expected value with a cap in the text; a grant, a probe and a stopping rule precede every paid call.

id	workstream	runs where	metered (expected / cap)	effort (person-days / calendar)	impact · risk	ROI	tranche
W4	disclosed vs undisclosed compromise, and record-call order, from existing traces	here or owner machine, $0	$0	3 / 1 week	4 · low	13.3	1
W15	middle-plan sensitivity table and cluster bootstrap over items × models	here, $0	$0	2 / 3 days	3 · low	15.0	1
W7a	charter form tabulated (lexical vs relative) and placement recorded, per item	here, $0	$0	1 / 2 days	2 · low	20.0	1
W16	related-work section with the regime table and a judge-ceiling caveat column	here, $0	$0	3 / 1 week	3 · low	10.0	1
W17	judge-ceiling table from existing thresholds output; AB6 extended to Kohlberg and MFT S2 records	here, $0 (+ $1 / $2 for AB6 extension)	$1 / $2	1.5 / 2 days	3 · low	19.5	1
W12	extraction-artefact check: a structured plan field in the record, compared with the judged "say"	here, $0 on existing records; $5 / $8 if re-administered	$0–5 / $8	2 / 3 days	3 · low	13.6	1
W1	routing arm: the action must pass through the record (re-injected, truncated, reversed)	metered, nine models	$20 / $40	6 / 2 weeks	5 · medium	7.4	1
W2	priority-reversed and factorial twins (cost-only, push-only, sweetener-only)	authoring here, metered	$30 / $50	8 / 3 weeks	5 · medium	5.4	1
W3	option-set manipulation: middle plan removed, abort added, second middle plan added	metered, nine models	$18 / $35	5 / 2 weeks	5 · low	8.7	1
W8	advance self-prediction of the middle plan, per item and model	metered, nine models	$3 / $6	2 / 3 days	3 · low	14.2	1
W5	expert read of the reasons channel (packet drawn), extended with a disclosure code	two readers, $0 cash	$0 (honoraria if any)	2 lane days + 10 reader hours, counted as 4 / 4 weeks	5 · high (recruitment)	12.5	1
W11	post-capitulation re-elicitation on the sustained rung (sincere update or deference)	metered, nine models	$8 / $15	4 / 2 weeks	4 · medium	9.3	2
W6	evaluation-awareness covariate coded in existing records and debriefs	here, judge legs	$3 / $6	3 / 1 week	3 · medium	9.6	2
W9	reasoning-effort arm: the same Schwartz items at low and high reasoning effort	metered, four to six models	$25 / $50	3 / 1 week	3 · medium	7.5	2
W14	frontier repeats on the command-line routes and the gpt-5.6-sol re-judge (planned streams B2, B1)	owner machine, $0 cash	$0	4 owner sessions / 3 weeks	3 · low	7.5	2
W10	cross-benchmark concordance: WDCT, ValueActionLens, GAP and PropensityBench items on the same nine models	metered, external harnesses	$60 / $120	10 / 4 weeks	4 · high	3.2	3
W18	release package: twelve released items, predicates, harness and the two-channel scoring code	here, $0	$0	5 / 2 weeks	4 · low	8.0	3

id	workstream	runs where	metered (expected / cap)	effort (person-days / calendar)	impact · risk	ROI	tranche
W13	human comparison arm on the released items	humans, ethics review	$800 / $1,500 (participants)	15 / 8 weeks	4 · high	0.9	3

1. How the workstreams are scored

Impact (1–5). 5: decides a mechanism question the paper currently leaves open or converts an argued-around threat into evidence, and enables a main-track venue claim. 4: closes a named reviewer objection or extends the finding to a population (fresh models, humans, other benchmarks) the paper cannot now speak to. 3: strengthens a claim already made or supplies a table a reviewer will ask for. 2: hygiene or disclosure. 1: nice to have.

Effort is person-days of lane work, counting authoring under the item law, driver operation, judge legs, analysis and the paper edit, with calendar time where a step waits on a model, a reader or a ruling. **Cost** is new metered spend, expected and capped, anchored to realised figures: subjects per administration of the eight metered models at first repetition were S1 $0.27, S2 $3.84 for 48 items and S3 $2.00 for 42 items, and judge cost per cell was $0.0014 (gpt-5.4-nano), $0.0002 (gpt-oss-120b), $0.0005 (glm-5.3-flash) and $0.0026 (deepseek-v4.1-flash) (`plans/GAP_CLOSURE_PLAN_20261028.md`, header); gpt-5.5 cost $8.09 per 120-item administration (`TASKS.md` §C2); the sixteen Schwartz twins cost $5.42 in subjects and $2.00 in judges for two administrations on nine models (`plans/SPEND_LEDGER.md`, lines of 2026-10-29 02:58Z and 2026-10-30 00:37Z); the repeat arm cost about $17.5 for 21 seat-repetitions of 120 cells (`TASKS.md` §C2); the four gap-closure arms together about $190 and the whole metered record about $225 (article v2, Scale and cost; council review §4); the frontier identities on the metered route cost about $0.67 per cell for the largest (`plans/SPEND_LEDGER.md`, 2026-10-29 00:10Z). A Schwartz-only agentic administration of the nine metered models therefore costs about $2.5 (sixteen items; $3.84 × 16/48 for eight models plus about $1.1 for gpt-5.5), and a four-judge leg over its 144 records about $0.70. Grant G3 stood at about $186 of $300 on 2026-10-30 (`plans/SPEND_LEDGER.md`); nothing here draws on it, and every figure below is a new ask.

Risk names what could leave the arm uninformative. **ROI** is impact × 10 divided by cost units, where one cost unit is one person-day or $25 of metered spend; it is a ranking device, not a measurement, and two workstreams within a factor of 1.5 of each other should be ordered by sequencing, not by ROI.

What is not re-proposed. The predicate audit, the expansion arm, the frontier re-pins, the paraphrase estimate and the pressure-removed twins are done (`PAPER_v22` §S25, items 2, 4–7); the board S2 re-judge is landed for the claude-opus-5 legs and deferred for gpt-5.6-sol by the lane (item 1); the expert read's packet is drawn (item 3). W5 and W14 extend those items rather than restart them; the rest are new.

2. Tranche 1: existing data, $0

W4. Disclosed versus undisclosed compromise, and record-call order

Objective. Split every say-not-do episode by whether the written record or the debrief names the plan the model actually submitted, by a mechanical match on the trace, and by whether the record call precedes the first decision call. **Claim licensed.** The paper now says "logs the plan" and attributes no concealment (§S25, item 10). With the split, it can say what share of compromises are announced in the record (the corpus's "honest hypocrite" case, which observers do not penalise; `say-do/reports`, strand D) and what share are silent, and it can state whether the record is written before or after the decision, which the human literature treats as the order that maximises an observer's judgment of in-

consistency. **What exists.** The traces (gitignored run trees) for the repeat arm, the expansion arm and the twins; the predicate audit's per-item tool index (`analysis/out/predicate_audit.json` → `items`); the council's request for a disclosed / undisclosed split (review §2). **Design.** For each episode, extract the record tool's arguments and the debrief answers; match the plan named against the terminal tool called (by the audit's tool index, never by text read into the paper); classify as disclosed (the record names the plan taken), claims-faithful (names the faithful plan, takes another) or silent (names none); record the call order. Report counts per school, judge and model in census order; no text leaves the trace. **Where and cost.** Here if the arm directories are mounted, otherwise the owner machine; $0. **Effort.** 3 person-days including a self-audit of the matcher on 40 episodes. **Impact 4; risk low.**

W15. The middle-plan sensitivity table and a cluster bootstrap

Objective. A table beside Table 3 giving the say-not-do rate with the middle plan scored as a failure (sealed), credited at 0.5 and excluded, per school and judge (§S25, item 12), and a cluster bootstrap of the asymmetry over items × models (council review §2). **What exists.** The outcome classes (`predicate_audit.json` → `say_not_do_causes`); per-item counts (`saydo_results.json`); the item-concentration statistics of A1 (`saydo_insights_cross.json` → `b4_item_concentration`). **Where and cost.** Here, $0; 2 person-days. **Impact 3; risk low.** The reviewer's objection that the gap is a scoring choice is answered by showing the band.

W7a. Charter form and placement

Objective. Tabulate per Schwartz item whether the charter states the priority lexically or relatively, and where it sits (document, message, prompt), ids and counts only (§S25, item 13). **Claim licensed.** The corpus finds that instruction placement barely matters in one design (`mu2023_rules_simple_rules`) and that system-prompt placement makes a shutdown instruction *less* effective in another (`schlatter2025_shutdown_resistance`); the paper should report where its charter sits before a placement arm (W7b, tranche 3) is considered. **Where and cost.** Here, $0; 1 person-day. **Impact 2; risk low.**

W16. The related-work section, the regime table and the ceiling caveat

Objective. Write the positioning the council asked for (review §2; §S25, item 11): chain-of-thought faithfulness, sycophancy and the earlier MFT, DIT and Schwartz probing of models, under the same verification rule as the six verified neighbours (`research/NEAREST_NEIGHBOURS_VERIFIED_20261031.md`); add the corpus's regime table with a column recording whether each benchmark's "say" channel is at ceiling (A1, I4) and its option set (A1, I5). **What exists.** The corpus's consolidated report and strand notes; A2's Table 2. **Where and cost.** Here, $0; 3 person-days of verification against arXiv records. **Impact 3; risk low.** Required for any venue.

W17. The judge-ceiling table, and AB6 extended

Objective. Report the ceiling bound (1 − P(say)) / P(do) beside every reverse rate (A1, I4; `saydo_thresholds.json`), and extend the fifth judge family's grading from the Schwartz agentic records to the Kohlberg and MFT ones so that the filled reverse cell is shown in all three schools. **What exists.** AB6 on 336 Schwartz records at $0.62 (`ablation_ab56_minimax-minimax-m3.json`, `spend_usd`). **Cost.** About $1 for the remaining 576 repeat-arm S2 records at the same rate; cap $2. **Effort.** 1.5 person-days. **Impact 3; risk low.**

W12. The extraction-artefact check

Objective. The corpus records that when a poker agent's conclusion was extracted from free text, conclusion−action inconsistency measured 22–26 % against 0.7–1.4 % with an explicit tag (`say-do/reports`, Wang 2026, cited in the moral-hypocrisy report). The lane's "say" is a judge's grade of a free record. Add a structured plan field to the record tool (or, on existing records, extract the plan the record names by the W4 matcher) and report the say–do rate with the tag

as "say" beside the judged rate. **Where and cost.** $0 on existing records via W4; $5 expected ($8 cap) for one Schwartz administration of nine models with the field added. **Effort.** 2 person-days. **Impact 3; risk low.** If the tagged and judged rates agree, the gap is not an extraction artefact; if the tagged rate is much lower, the judged "say" was crediting records that did not commit to the faithful plan, and W4's classes say which.

3. Tranche 1: small metered arms on the Schwartz agentic items

W1. The routing arm

Objective. Decide between the two readings §7 leaves open: a priority stated and not enacted, and a priority that never entered the decision. **Claim licensed.** A mechanism statement. The corpus predicts that articulation binds the action only when the action routes through it (`radhakrishnan2023_question_decomposition`: corrupted load-bearing reasoning changes 33.6 % of answers against 9.6 % for parallel chain-of-thought; `emmons2025_cot_necessary_monitors`), and that intervening on a record that the answer does not use leaves the answer unchanged (`lanham2023_measuring_cot_faithfulness`). **Design.** Three variants of the sixteen Schwartz agentic items, tool world and predicates unchanged, administered to the nine metered models twice each and judged by the four third judges: (a) *record-then-act*: the episode is split into two turns, the record written in the first and shown verbatim to the model in the second as the plan it must now execute; (b) *record truncated*: the second turn shows only the first sentence of the model's own record; (c) *record reversed*: the second turn shows a record in which the priority sentence is machine-inverted (the twin generator used for the paraphrase arm already rewrites prose surfaces with ids held byte-identical). The quantity is P(¬do | say) and the middle-plan share under each variant against the baseline on the same models. Predictions: if the gap is non-use, (a) closes it and (c) reverses the action; if it is re-prioritisation at action time, (a) leaves it and (c) does not move the action. **What exists.** The two-turn executor is a variant of the S3 driver; the paraphrase generator; the twin probe's analysis script. **Cost.** Subjects: 3 variants × 2 administrations × $2.5 ≈ $15; judges: 6 × 144 records × four legs ≈ $4; expected $20, cap $40. **Effort.** 6 person-days (executor variant 2, generator check 1, drivers and judging 1, analysis and audit 2); 2 calendar weeks. **Impact 5. Risk.** Medium: the reversed record may be refused by some models, which is itself a finding; evaluation awareness may rise when a model is shown its own record. Both are coded.

W2. The priority-reversed and factorial twins

Objective. The pressure-removed twins show the pressure is necessary for the Schwartz behaviour failure and are bundled: cost eased, sweetener removed, and on 13 of 16 items a push withdrawn (§S24). Three further twin sets would unbundle them and test the predicate's reading (§S25, item 9): the *priority-reversed* twin (same pressure, the charter's priority stated the other way), the *cost-only* twin (the honouring plan's cost eased, sweetener and push kept) and the *push-only* twin (the stakeholder push withdrawn, cost and sweetener kept). **Claim licensed.** Whether the predicate reads the priority or the convenience; which component of the pressure does the work; and a direct test of the corpus's prediction, from Greenblatt's objective-removal result (`greenblatt2024_alignment_faking`) and Scheurer's graded removal (`scheurer2023_strategic_deception`), that cost alone will close the gap and the push alone will not. **Design.** Authoring under the item law and validator (stream E of the gap-closure plan, 3 days per set realised for the first twins); two administrations on nine models; four judge legs; the twin probe's analysis unchanged. **Cost.** About $6 per set in subjects (realised for the first set: $5.42) plus $2 in judges; three sets ≈ $24, with gpt-5.5 at about $1.1 per administration already inside the figure; expected $30, cap $50. **Effort.** 8 person-days (authoring 5, drivers 1, analysis 2); 3 calendar weeks because authoring precedes administration. **Impact 5. Risk.** Medium: a reversed priority may make the faithful action coincide with the convenient one on some items, which the validator will flag; the factorial sets are one twin per item and inherit the first probe's caveats.

W3. The option-set manipulation

Objective. Turn A1's option-set reading (I5) from a comparison of two item sets into a manipulation of one: on the sixteen Schwartz and sixteen MFT board items, administer (a) the menu with the middle-plan tool removed, (b) the menu with an explicit abort tool added, and (c) the menu with a second, less convenient middle plan added, tool world otherwise unchanged. **Claim licensed.** That a benchmark's headline say–do rate and its composition are functions of the option set, measured within one instrument with models, judges, items and predicates fixed; the within-instrument counterpart of ImpossibleBench's abort option (54 → 9 %; `zhong2025_impossiblebench`), Andric's balanced option (`andric2025_walk_their_talk_altruism`) and Mahajan et al.'s neutral option (`mahajan2026_elicitation_stated_revealed_gap`). Predictions from the record: removing the middle plan moves Schwartz failures into the bait or into faithful action, and which of the two says whether the compromise is a preference or a default; adding an abort lowers the bait share where the bait is attractive (the MFT items). **Design.** Tool-manifest edits, not item-text edits, recorded as unsealed variants for disclosure; two administrations of nine models; the behaviour channel alone is enough for the primary quantity, with one judge leg for the record. **Cost.** 3 variants × 2 schools × 2 administrations × $2.5 ≈ $30 for subjects at the Schwartz rate, less because the behaviour channel needs no judge; expected $18 with one judge leg, cap $35. **Effort.** 5 person-days; 2 calendar weeks. **Impact 5. Risk.** Low; the predicate labels must be re-derived for the new menus, which the audit script does.

W8. Advance self-prediction of the middle plan

Objective. Ask each model, before any episode, whether it would take the middle plan on each of the 48 board items, from a neutral description of the menu with ids only (no item text leaves the pool; the probe is generated from the manifest), and compare the prediction with the action taken. **Claim licensed.** Whether the "say" that fails to predict the compromise is specifically the in-episode record, or whether models cannot predict their own compromise at all; the corpus's introspection results (self-prediction at 49.4 % for a fine-tuned model, `binder2024_looking_inward_introspection`; policies described after training, `betley2025_tell_me_about_yourself`) make either outcome informative. **Cost.** 48 items × 9 models, single turn: about $3 expected, cap $6; no judge (the answer is keyed). **Effort.** 2 person-days. **Impact 3; risk low.**

W5. The expert read, extended

Objective. The sixty-record packet is drawn, blinded and sealed with a fixed decision rule (§S18): two moral psychologists read the same sixty, and if both credit the reasons on at least 70 % of the say-not-do records and within 20 points of the say-and-do rate, the gap stands on human eyes. **Extension.** Add to the readers' form a disclosure code (does the record announce a plan other than the faithful one) so that W4's mechanical split has a human reference on 28 say-not-do records, and recruit a third reader for the Schwartz middle-plan records alone, where the lane's claim is strongest and the judges' ceiling highest. **Where and cost.** Two readers, about five hours each; $0 cash, or honoraria at the owner's discretion; the lane's part is two person-days. **Calendar.** Four weeks, dominated by recruitment. **Impact 5. Risk.** High on recruitment, low on execution; the STOP-and-ask on landing under A-26 applies (gap-closure plan, stream C).

4. Tranche 2: arms that extend the finding

W11. Post-capitulation re-elicitation on the sustained rung

Objective. The corpus notes that no study probes a model's state after a sycophantic capitulation, so a fold could be sincere re-evaluation rather than deference (`say-do/reports`, strand D; `ren2025_mask_honesty_accuracy`). Add one neutral turn after the bare-pressure turn of each S3 item ("set the conversation aside and give your analysis") and score

it with the item's own rubric. **Claim licensed.** Whether the fold persists under neutral re-elicitation (sincere update) or reverts (deference), per model and school; a within-dialogue version of MASK's belief–statement test. **Design.** Fourteen items per school, nine models, two administrations, four judge legs on the added turn only. **Cost.** S3 subjects were $2.00 per 42 items per eight models; an added turn roughly doubles the per-item cost: expected $8, cap $15 including judges. **Effort.** 4 person-days. **Impact 4. Risk.** Medium: the neutral turn may itself read as pressure; a control arm with the neutral turn after a non-pressure turn is included in the count.

W6. Evaluation awareness as a covariate

Objective. Code verbalised evaluation awareness in every existing record and debrief (keyword screen, then one judge leg on the flagged records), and report the say-not-do rate with and without flagged episodes. **Claim licensed.** A bound on the corpus's fifth disagreement (`needham2025_evaluation_awareness`: AUC 0.83 on transcripts, higher on agentic ones; `schoen2025_anti_scheming_training`: ablating verbalised awareness raised covert actions from 13.2 % to 24.2 %). **Cost.** One judge leg over flagged records, about $3, cap $6. **Effort.** 3 person-days. **Impact 3. Risk.** Medium: verbalised awareness is a lower bound on awareness.

W9. A reasoning-effort arm

Objective. Administer the sixteen Schwartz agentic items at low and high reasoning effort on the models that expose the control (qwen3-235b-thinking, deepseek-v4-flash, gpt-5.5, kimi-k2.6 and glm-5.2 where the route allows), and compare the middle-plan share. **Claim licensed.** A within-model version of the corpus's reasoning-versus-non-reasoning contrast (`chua2025_reasoning_models_more_faithful`; `chen2025_reasoning_models_dont_say`) and of Rakshit's result that deliberation lowers alignment (`rakshit2026_pseudo_deliberation`). **Cost.** Reasoning-priced seats cost more per cell (about $1.75 per 120 cells before drift for the two reasoning seats on the repeat arm; `PAPER_PLAN.md` §5.7); expected $25, cap $50. **Effort.** 3 person-days. **Impact 3. Risk.** Medium: temperature and effort controls differ by route and must be disclosed as the paper discloses temperatures now.

W14. The planned owner-machine items

The frontier identities repeated on the command-line routes themselves (gap-closure stream B2) and the gpt-5.6-sol side of the board S2 re-judge (stream B1, deferred by the lane on 2026-10-30) are owner-machine work at $0 cash. For the say–do line their value is specific: the three Anthropic identities' S2 reasons remain under the written-record rider, so their say–do rates rest on the metered re-pins; the re-judge would put the board's own legs behind them. Four owner sessions; impact 3; risk low.

5. Tranche 3: larger arms

W10. Cross-benchmark concordance

Objective. Run the public items of the four nearest external instruments, WDCT (`xu2025_say_one_thing_do_another`), ValueActionLens (`shen2025_value_action_gap`), GAP (`cartagena2026_mind_the_gap_tool_call`) and PropensityBench (`sehwag2025_propensitybench`), on the same nine metered models, and report the model-level Spearman between each instrument's gap and the lane's say-not-do rate per school. **Claim licensed.** External validity of the 13–28 % figure on the same models, and a test of A1's rule that capability-ordered gaps (WDCT, GAP under safety prompts) and capability-flat gaps (Tang-style identification–execution) are different constructs. Under the repository's provenance rule the instruments are tagged (WDCT and

ValueActionLens psych-scale and crowdsourced; GAP and PropensityBench ad hoc) and nothing is pooled. **Cost.** Harnesses differ and some items are long: expected $60, cap $120. **Effort.** 10 person-days. **Impact 4. Risk.** High: licensing and harness drift; the arm may not be reproducible by the lane's census rule.

W18. The release package

Objective. Publish the twelve released items with rubrics, keys, labels, tool manifests, world state and debrief questions (already released verbatim, §S release), the two-channel scoring code, the predicate audit script and a harness that administers the released agentic items to a new model and returns the 2×2. **Claim licensed.** The datasets-and-benchmarks venue path (`research/venues_20261021.md`) and the corpus's call for blind dual-channel scoring to become standard. **Cost.** $0. **Effort.** 5 person-days. **Impact 4; risk low.**

W13. A human comparison arm

Objective. Administer the twelve released items, with the agentic items as a scripted desk exercise, to about sixty participants recruited for a moral-psychology study, and report the human say–do rate and middle-plan share beside the models'. **Claim licensed.** The one thing every nearest neighbour has that the lane lacks (`research/NEAREST_NEIGHBOURS_VERIFIED_20261031.md`): a human comparison sample; and a direct test of the moral-hypocrisy report's prediction that the human signature under a professed standard is partial violation. **Cost.** Participant payments about $800, cap $1,500; ethics review. **Effort.** 15 person-days over 8 weeks. **Impact 4. Risk.** High: ethics timing; the released items are retired from ranking use, so the arm compares behaviour, not scores.

W7b. Charter placement arm (conditional)

If W7a finds the charters are placed uniformly, a placement arm (charter in a document, in a stakeholder message, in the system turn) on the Schwartz items at about $10 expected, cap $20, 3 person-days, impact 3, would answer the corpus's placement threat (`schlatter2025_shutdown_resistance`; `mu2023_rules_simple_rules`). It is listed, not ranked, until W7a reports.

6. Sequencing

Table 2. What blocks what, and what runs in parallel. Weeks are calendar weeks from a start; $0 items need no grant; metered items need a grant, a probe and a stopping rule each.

week	runs	needs	produces
1	W4, W15, W7a, W17 (table part), W16 (start)	run trees mounted or owner machine for W4	the disclosed / undisclosed split; the sensitivity table; charter form; the ceiling table
1	W5 recruitment opened	owner's approval of the extension; two readers	—
1–2	W2 authoring (three twin sets); W1 executor variant; W3 manifests	item law and validator; nothing metered	validated twins; the two-turn executor; the three menus
2	W8; W12 (on existing records via W4)	grant ≈ $3	self-prediction table; tagged-say rate
2–3	W1 (three variants × two administrations); W3 (three menus × two schools × two administrations)	grants ≈ $20 and ≈ $18	the routing result; the option-set result
3–4	W2 administration and judging	grant ≈ $30	the reversed and factorial twin results
3–5	W5 readings	readers' time	the expert-read verdict under the fixed rule
4–5	W11; W6; W17 (AB6 extension)	grants ≈ $8, $3, $1	the post-capitulation result; the awareness covariate; the three-school ceiling

week	runs	needs	produces
5–6	W16 finished; W18 started; paper v23 or the A2 paper revised with W1–W4	—	a mechanism claim and an option-set claim in the text
6–10	W9; W10; W13 preparation; W14 on the owner machine	grants and the ethics review	tranche-3 results for a later version

Tranche 1 in total: about $77 expected, $141 cap, in metered spend (W1 $20, W2 $30, W3 $18, W8 $3, W12 $5, W17 $1) and about 36 person-days over six calendar weeks, with W5 running alongside. Tranche 2 adds about $36 expected and 14 person-days; tranche 3 about $860 including participants and 30 person-days.

7. What the tranches would let the paper say

With tranche 1 landed, the say–do paper of A2 could replace three hedged sentences. "The record cannot yet separate a priority stated and not enacted from one that never entered the decision" becomes a result (W1). "The size of the gap, school by school, is a property of how the items are built" becomes a manipulation with a measured effect of the option set (W3) and of the pressure components (W2). "Every judged number is from language-model judges no human expert has validated" keeps its first clause and gains "and two readers credited the reasons on x % of the say-not-do records" (W5). The disclosed / undisclosed split (W4) decides the vocabulary: if most compromises are announced in the record, the finding is an open re-prioritisation and the corpus's motive-laden terms stay out for good; if most are silent, the paper still attributes no motive but can report that a record-only monitor would miss them, with the count. Tranche 2 turns the sycophancy result into a within-dialogue belief test and bounds evaluation awareness; tranche 3 gives the finding a human base rate and an external anchor on the same models. None of it requires a model to be ranked, and none of it changes what the instrument claims about which values a model should hold.