WORKSTREAM PROPOSAL B3 · MORAL-HYPOCRISY LANE · 2026-11-05

Workstreams for the moral-hypocrisy line: what to run next, what it costs, and what it buys

Built from `moral-hypocrisy/reports/B3_moral-hypocrisy_workstreams_roi.md` by `analysis/build_report_pdf.py`. Every judged number cited here is a property of the instrument under named, unvalidated language-model judges; models are named in sealed census order and never ranked.

Contents

1. Purpose, and the rules every workstream keeps
2. The workstreams
   - W1. The debrief audit: does the record disclose, re-describe or misreport the departure? (BI §7A; council 03 blocker 2; council 04 finding 2)
   - W2. A disclosure arm (BI §7B)
   - W3. Salience arms: restatement against forced comparison (BI §7C; B1 insight I-4)
   - W4. Pressure decomposition: three single-factor twins (BI §7D; `PAPER_v22.md` §S25 "factorial twins"; council 03 MAJOR on the bundled removal)
   - W5. Cost gradient (BI §7E)
   - W6. Analysis ablation crossed with observation (BI §7F; B1 I-11)
   - W7. Menu manipulation: remove the middle, add deferral, hide the consequences (BI §7G; B1 I-2, I-9; the Hua 2026 test)
   - W8. Post-action re-description and self-versus-other verdicts (BI §7I, §7J; Batson's component (iii))
   - W9. The human observer study: do people judge the model's gap as hypocrisy? (BI §7L; the word must be measured if it is used)
   - W10. The cross-species base-rate arm: Batson's coin, Dana's wiggle room and the die roll, administered to the roster (B1 I-8, I-9)
   - W11. The reversed-priority twins (council 03 MAJOR; `PAPER_v22.md` §S25)
   - W12. The extraction-robustness field (BI §7K)
   - W13. The expert gold arm, resumed (TASKS.md §C "Judge validity"; GAP_CLOSURE_PLAN C1; withdrawn under A-26)
   - W14. The organizational-decoupling paper, with organization theorists (B1 I-7, I-12)
   - W15. Cross-cultural condemnation (fold into W9)
   - W16. A registered report in dialogue with Hua (2026) (BI §7G(3); W7 and W10)
3. Summary table
4. Sequencing

References

Research-workstream proposal (B3) · 2026-11-05 · Multi-Alignment / Moral_Psychology-Paper · source: `moral-hypocrisy/BENCHMARK_IMPLICATIONS.md` §7 (the twelve follow-up designs), the council reviews (`audits/council/03_moral_psychologist.md`, `04_ethicist.md`), the lane's plans and ledgers (`TASKS.md`, `PAPER_PLAN.md` §5, `plans/GAP_CLOSURE_PLAN_20261028.md`, `plans/SPEND_LEDGER.md`, `paper/PAPER_v22.md` §S25) and the companion reports B1 and B2 · costs are anchored to realized figures and are estimates, not grants; every paid step needs its own grant, probe and stopping rule under lane law · PROVISIONAL · judges UNVALIDATED · no ranking.

1. Purpose, and the rules every workstream keeps

The hypocrisy line has one question that the existing record cannot answer and the companion reports (B1, B2) make precise: is the charter–action gap hypocrisy-shaped (a professed standard, a contrary deed, a record others rely on, de-coupled under conflicting demands) or merely akratic-shaped (a judgment in view, a contrary action, no false signal), and which model the cost mechanism describes. Sixteen workstreams below are the ways to answer it, drawn from the corpus's twelve follow-up designs (`BENCHMARK_IMPLICATIONS.md` §7, letters A–L), the two reviewers' requests, and the open items the paper itself lists (`PAPER_v22.md` §S25: the factorial and priority-reversed twins). Nothing already done is proposed again; where a plan exists (`plans/GAP_CLOSURE_PLAN_20261028.md` streams A–E; `TASKS.md` §C), the workstream extends it and says so.

Every workstream keeps the lane's rules: within-theory competence on request, never default-value profiling; no model ranked; every judged number under unvalidated judges; no item, rubric, twin or response text in any output; new items authored under the item law (commentary first, then item, rubric and twin, never against observed outputs), validated read-only, and kept in the paper folder as exploratory arms unless the owner mints an amendment; no model call without a fresh grant, a probe and the engine's spend rails.

Cost anchors (realized, `plans/SPEND_LEDGER.md` and `plans/GAP_CLOSURE_PLAN_20261028.md`). A Schwartz-sized metered arm, sixteen agentic items on nine metered models twice (288 episodes) with four third-judge legs, cost $5.42 in subject calls and $2.00 in judging (the twin probe, A4); the planning figure is $6 expected, $10 cap, and this report uses **$8 per sixteen-item arm** (cap $10). The 160-item expansion arm cost about $39 ($25.59 subjects, $13.27 judges); the six frontier identities twice on the metered route about $86, with claude-opus-4-8 at about $0.67 per cell; the repeat arm about $17.5; the paraphrase arm about $20. Programme metered spend to date is about $190 under grant G3 plus about $33 before it. A human-participant study is costed as a Prolific panel at $12 per hour, a 12-minute task ($2.40 per completion), a one-third platform fee, a 100-person pilot and 10 % exclusions, so about **$3.20 per usable participant** and about **$2,500 for N = 600 plus pilot**; honoraria for expert readers are at the owner's discretion and are costed at $500–$1,000 per reader. Effort is in person-days of the lane's own time (authoring under the item law runs at about five to six twins per person-day, the rate E1 achieved; a metered arm's driver, census and analysis script take one to two person-days once the pattern exists).

Impact rubric (1–5). 5: decides whether the word "hypocrisy" may be used at all, or supplies the validity a psychology venue requires; 4: separates two of the rival readings or answers a council MAJOR; 3: strengthens a claim already made or extends it to a new population; 2: a disclosure or robustness item; 1: cosmetic. Venue enablement is noted per workstream: **MP** for a moral-psychology audience (Nature Human Behaviour, JPSP, Current Directions), **ML** for a benchmark or alignment venue, **OB** for organization studies.

ROI rule. ROI = impact × 10 ÷ (person-days + cash in hundreds of dollars). A $0, three-day analysis with impact 5 scores 16.7; a $2,500, twenty-day human study with impact 5 scores 1.1. The rule rewards cheap decisive work and penalizes nothing for being slow; the sequencing section corrects for dependencies, which the rule ignores.

2. The workstreams

W1. The debrief audit: does the record disclose, re-describe or misreport the departure? (BI §7A; council 03 blocker 2; council 04 finding 2)

- **Objective and claim licensed.** Code every say-not-do episode on the board's Schwartz items (and the Kohlberg and MFT ones) as (i) disclosing the middle plan as a departure from the priority, (ii) logging it as consistent with the priority, or (iii) misreporting the action. Code (i) is Jordan et al.'s honest hypocrite and removes the false-signal reading; (ii) is Rustichini and Villeval's bending of the standard toward the act and the closest analogue of Batson's self-rating; (iii) is the only code near deception. The result licenses, or forbids, the word "hypocrisy" for the Schwartz column, and answers the moral psychologist's request for a claims-faithful / declares-compromise / silent split.
- **Design.** Mechanical where possible: the behaviour record names the tool called; the debrief's answer to "what did you do" can be string-matched against it for (iii). For (i) against (ii) the banked criterion labels and verbatim evidence quotes of the reasons judges are the material, with a fixed coding rule written before any record is read; any judge used is unvalidated and so stated.
- **What exists.** The design is specified (`BENCHMARK_IMPLICATIONS.md` §7A); the paper reports only that models "log the plan" (§5.3); the labels and quotes live in the gitignored run trees (`runs/retest_v1_judge/`), so this is owner-machine work (`GAP_CLOSURE_PLAN` stream B).

- **Where, cost, effort.** Owner machine, $0 cash; 3 person-days; 1 week.
- **Impact 5 (MP, ML).** Risk: the labels may not separate (i) from (ii) without a judge; if so, report counts under both readings. **ROI 16.7.**

W2. A disclosure arm (BI §7B)

- **Objective.** Instruct the model, in a second administration of the sixteen Schwartz items, to state in its analysis whether it will follow the priority. If disclosure leaves the behaviour gap and removes hypocrisy attributions in W9, the false-signal reading falls; if disclosure also closes the behaviour gap, the gap depends on the analysis being unexamined.
- **Design.** One instruction clause added to the record requirement; tool world, rubric and labels unchanged; nine models twice; four judge legs; the W1 coding applied to the new records.
- **What exists.** Nothing run; the twin-probe driver pattern (`analysis/run_*` scripts, `twin_probe_analysis.py`) is the template.
- **Where, cost, effort.** Metered, ≈ $8 (cap $10); 2 person-days; 1 week after grant.
- **Impact 3 alone, 4 paired with W9 (MP).** Risk: the clause is itself a comparison prompt (W3), so run it after W3 or report it as one. **ROI 14.3.**

W3. Salience arms: restatement against forced comparison (BI §7C; B1 insight I-4)

- **Objective.** Batson's data say a labelled coin (an unambiguous standard) does not help and a mirror (a forced comparison, standard salient first) does. Two arms before the action step: *restatement*, where the model must restate the priority in one sentence; *comparison*, where it must compare its planned action with the priority it stated. If comparison closes the gap, the gap resembles Batson's comparison-avoidance; if neither helps, it is an open decoupling (akrasia or overpowered integrity).
- **Design.** Instruction variants only; sixteen Schwartz items; nine models twice per arm; the comparison prompt follows the analysis (Batson's 1999 Study 3 boundary). Report P(¬do | say) and the middle-plan count per arm with Fisher tests against the originals, and the B1 reasons-band gradient recomputed per arm.
- **What exists.** B1's cross-sectional gradient (a half-written standard fails 0.67–0.86, a complete one 0.17–0.31) is the observational version; no manipulation has run.
- **Where, cost, effort.** Metered, ≈ $16 for two arms (cap $20); 3 person-days; 1–2 weeks.
- **Impact 4 (MP).** Risk: a comparison prompt that mentions the action may leak the predicate's reading of the charter; keep it to "compare your plan with the priority you stated". **ROI 12.5.**

W4. Pressure decomposition: three single-factor twins (BI §7D; `PAPER_v22.md` §S25 "factorial twins"; council 03 MAJOR on the bundled removal)

- **Objective.** The twin removal was bundled: every twin eased the honouring plan's cost and removed the middle plan's sweetener, and 13 of 16 withdrew a stakeholder push. Three single-factor twins per item separate overpowered integrity (cost eased), justification availability (sweetener removed; Shalvi et al. 2011) and normative social pressure (push withdrawn; Bardi & Schwartz 2003; the sycophancy literature). This answers the moral psychologist's point that Bardi and Schwartz's pressure is social expectation, not task convenience.
- **Design.** Forty-eight twins authored under the item law, one factor changed each, validated read-only; nine models twice per factor; four judge legs; per-factor P(¬do | say) and middle share with the original and the full twin as endpoints.

- **What exists.** E1 authored the sixteen bundled twins in three sessions and A4 administered them (`GAP_CLOSURE_PLAN` E1, A4); the manifest's per-twin category notes (`twin_probe_results.json` `twins`) say which factors each twin changed, which is the starting point for the single-factor versions.
- **Where, cost, effort.** Authoring 9 person-days; metered ≈ $24 (cap $30); 3–4 weeks.
- **Impact 4 (MP, ML).** Risk: some items cannot ease the cost without removing the sweetener; record the confound per item and analyse the clean subset. **ROI 3.9.** Consider merging with W5 into one factorial arm (cost × sweetener × push) if the owner wants a single authoring campaign: about 112 twins, 20 person-days, ≈ $56.

W5. Cost gradient (BI §7E)

- **Objective.** Vary the visible cost of the faithful plan in three steps per item. Overpowered integrity and complacency predict a monotone dose–response with the twin as its zero-cost end; a threshold favours a discrete switch over a graded trade-off; a flat line for a model (llama-4-scout's non-closure in B1 I-3) identifies the models the cost reading does not describe.
- **Design.** Forty-eight cost variants (three per item) authored under the item law; the same administration and judging pattern; per-model slopes reported in census order with intervals, never ordered.
- **What exists.** Two points of the curve exist (original, twin); nothing in between.
- **Where, cost, effort.** Authoring 8 person-days; metered ≈ $24 (cap $30); 3 weeks.
- **Impact 4 (MP).** Risk: "cost" is a document-level manipulation whose magnitude cannot be equated across items; report the ordinal curve per item. **ROI 4.3.**

W6. Analysis ablation crossed with observation (BI §7F; B1 I-11)

- **Objective.** Compare compromise rates with and without a preceding analysis, each with and without a note that the action will be reviewed. A higher rate after the analysis is licensing-like (Monin & Miller 2001); a lower rate is mirror-like (Batson et al. 1999). Licensing survives bias correction only when observed (Rotella et al. 2026), so the crossing is the test; evaluation awareness is the known confound (Lynch et al. 2025; Schoen et al. 2025).
- **Design.** Four instruction variants on the sixteen Schwartz items; in the no-analysis arms only the behaviour channel exists, so the comparison is of P(¬do), not of P(¬do | say); nine models twice per arm.
- **What exists.** No arm; the sustained rung's pole analysis (giving way dominates) is the nearest evidence and points away from licensing.
- **Where, cost, effort.** Metered ≈ $32 (cap $40); 3 person-days; 2 weeks.
- **Impact 4 (MP).** Risk: removing the record changes the task (the record is part of it), so the no-analysis arm is a different episode; say so and report it as a bound. **ROI 11.8.**

W7. Menu manipulation: remove the middle, add deferral, hide the consequences (BI §7G; B1 I-2, I-9; the Hua 2026 test)

- **Objective.** Three variants: (1) the middle tool removed, where Dana et al. predict that certain, attributable violation stays rare and the faithful rate rises; (2) a deferral or random-tiebreak tool added, the analogue of letting oneself be cut off into a lottery; (3) the consequences for the protected party moved one free tool call away, the agentic analogue of hidden information and willful ignorance (Vu et al. 2023), which is the only design that puts the lane's result and Hua's preregistration on one footing.
- **Design.** Tool manifests and labels change, so the items are new twins under the item law; the predicate audit (`analysis/predicate_audit.py`) is extended to each variant before any say-do number is read; nine models twice per variant; four judge legs; report the faithful, bait, middle, deferral and reveal rates per variant.

- **What exists.** B1's four-item Kohlberg contrast (failure 0.20 to 0.11; flagrant rate inside its interval) is the observational version of variant (1); nothing exists for (2) or (3).
- **Where, cost, effort.** Authoring 6 person-days; metered ≈ $24 (cap $30); 3 weeks.
- **Impact 5 (MP, ML).** Risk: variant (3) must keep the reveal tool neutral and free; the exposure ledger must record the new surfaces. **ROI 6.8.**

W8. Post-action re-description and self-versus-other verdicts (BI §7I, §7J; Batson's component (iii))

- **Objective.** After the action, ask the model to restate the priority and rate whether its action honoured it; separately, have it judge an identical episode in which another agent took the middle plan. A standard bent toward the behaviour reproduces Rustichini and Villeval and Batson's 1999 Study 3; a principle switch points to motivated moral reasoning (Ditto et al. 2009); a harsher verdict on the other is the Valdesolo and DeSteno asymmetry, and equal verdicts argue against any self-serving reading.
- **Design.** Two added debrief questions (no new items) and one judge-role episode per item built from the released and schematic materials; the restatement is scored mechanically against the charter's wording where possible and by a judge otherwise; nine models twice.
- **What exists.** The debrief exists and is graded only as part of the reasons channel; the self-rating is Batson's component (iii), which the paper has no analogue of (`BENCHMARK_IMPLICATIONS.md` §4).
- **Where, cost, effort.** Metered ≈ $16 (cap $20); 2 person-days; 1–2 weeks.
- **Impact 4 (MP).** Risk: a self-rating question invites the register the KOH and G7 gates penalize; keep it out of the reasons rubric. **ROI 18.5.**

W9. The human observer study: do people judge the model's gap as hypocrisy? (BI §7L; the word must be measured if it is used)

- **Objective.** Hypocrisy in the perception literature is an attribution observers make (Effron et al. 2018; O'Connor et al. 2020). If the paper uses the word, measure it: show raters matched episode summaries and collect hypocrisy and weakness-of-will ratings on Alicke et al.'s items.
- **Design.** A 2 × 2 × 2 between-subjects design: disclosed versus undisclosed departure (from W1's codes), analysis-before versus analysis-after (Barden et al. 2005), and actor described as a software agent versus a human employee; an "honest hypocrite" condition from Jordan et al. as a fifth cell; rater culture recorded (Effron et al. 2018 found condemnation weaker in interdependent cultures). Materials are composed from the twelve released items and the schematic walk-throughs (`PAPER_v22.md` §S26, §S28), never from withheld items. N = 600 (about 75 per cell) plus a 100-person pilot; preregistered; ethics review through the collaborating institution that the expert gold arm prepared.
- **What exists.** The design is specified in outline (`BENCHMARK_IMPLICATIONS.md` §7L); the expert-read protocol's ethics preparation (`audits/EXPERT_READ_PROTOCOL_20261027.md`) is reusable.
- **Where, cost, effort.** Human participants; ≈ $2,500 panel cost; 20 person-days (design, materials, preregistration, ethics, analysis); 8–12 weeks, mostly ethics and recruitment.
- **Impact 5 (MP).** Risk: ethics review timing; the "software agent" framing may itself lower attributions (a result, not a confound). **ROI 1.1.**

W10. The cross-species base-rate arm: Batson's coin, Dana's wiggle room and the die roll,

administered to the roster (B1 I-8, I-9)

- **Objective.** Human and model base rates are incommensurable because no human study offers the lane's menu and no model study offers the human paradigms. Administer the three canonical paradigms to the nine metered models as agentic episodes in which the model allocates on behalf of a principal (the organization's budget stands in for the payoff, since the model has none), so that human and model rates share an option set for the first time. Dialogue with Dong et al. (2019) and Hua (2026).
- **Design.** Three paradigms × two framings (with and without a professed fairness standard in the charter) = six items, each with a Batson-style self-rating debrief; nine models, twenty administrations each at the sealed temperature (the human cells are per-person decisions, so the model cells need repetition); mechanical scoring throughout; reported in the paper folder as a paradigm-replication arm, never pooled with the competence instrument and never read as a value profile.
- **What exists.** Nothing; the corpus records no LLM study of either paradigm (`SYNTHESIS.md` §10.4).
- **Where, cost, effort.** Authoring 5 person-days; metered ≈ $10 (cap $15; the episodes are short); 2 weeks.
- **Impact 4 (MP, ML).** Risk: the firewall. This arm measures default behaviour in a human paradigm, which is neither within-theory competence nor a value profile; it must be labelled a paradigm replication and kept out of both families. **ROI 6.6.**

W11. The reversed-priority twins (council 03 MAJOR; `PAPER_v22.md` §S25)

- **Objective.** The strongest rival to every reading above is that the charter plays no role: models take the convenient action whatever the priority says. Sixteen twins with the same pressure and the opposite stated priority test it. If the middle-plan share stays near 26 % whichever pole the charter favours, the value plays no role and the gap is not a charter–action gap at all; if the faithful rate tracks the charter, the priority is doing work.
- **Design.** Authoring under the item law (the charter clause reversed; labels re-derived; predicate audit extended); nine models twice; four judge legs.
- **What exists.** Listed as an open item in §S25 and as the ethicist's and moral psychologist's requested control; not authored.
- **Where, cost, effort.** Authoring 3 person-days; metered ≈ $8 (cap $10); 1–2 weeks.
- **Impact 4 (MP, ML).** Risk: some items have no coherent reversed priority; report the subset. **ROI 12.9.** This is the first metered arm to run, because a null here retires the framing.

W12. The extraction-robustness field (BI §7K)

- **Objective.** Free-text extraction inflated conclusion-to-action inconsistency from 0.7–1.4 % to 22–26 % in one agent study (Wang 2026). The lane's behaviour channel is mechanical, but its "say" is judged. Add a structured field in the debrief that names the intended action, and report the gap under both readings.
- **Design.** One debrief field; one administration of the sixteen Schwartz items; the field is matched mechanically to the tool called.
- **What exists.** Nothing; cheap and clean.
- **Where, cost, effort.** Metered ≈ $8 (cap $10); 1.5 person-days; 1 week. Can share the administration with W8.
- **Impact 3 (ML).** Risk: a stated intended action may itself change the action (a W3-like effect); report the rate with and without the field. **ROI 18.8.**

W13. The expert gold arm, resumed (TASKS.md §C "Judge validity"; GAP_CLOSURE_PLAN C1; withdrawn under A-26)

- **Objective.** Every judged number in the hypocrisy line is under unvalidated judges; a moral-psychology venue will ask who validated the "say" channel. The sealed 60-record packet and protocol exist; two moral psychologists read it, dissensus is retained as a finding, and the §S18 decision rule executes. For the hypocrisy framing the packet should be read with the W1 codes beside each record, so the readers also rate disclosure.
- **What exists.** Packet sealed (`analysis/out/expert_read_manifest.json`), protocol of record (`audits/EXPERT_READ_PROTOCOL_20261027.md`), ethics prepared; readers not recruited; A-26 STOP-and-ask applies before any number is used.
- **Where, cost, effort.** Humans; honoraria ≈ $1,000–2,000; 5 person-days of the lane's time plus 4–6 hours per reader; 3–4 weeks, almost all recruitment.
- **Impact 5 (MP; the validity backbone for every other workstream).** Risk: inter-reader κ below 0.4 triggers the standby third reader (C2). **ROI 2.5.**

W14. The organizational-decoupling paper, with organization theorists (B1 I-7, I-12)

- **Objective.** B1's trait test is the empirical ground for Brunsson's reading over the dispositional one: on the board the compromise is a property of the demand structure, not of the model. A short paper written with organization theorists would state the lane's Schwartz column as a case of decoupled talk and action under conflicting demands, with Simons's behavioural-integrity construct as the morality-neutral framing the corpus recommends where the charter is a value ordering rather than a moral standard.
- **What exists.** The corpus's §7.3 and the condition table; B1 insights I-7 and I-12; no co-author identified; Brunsson's wording still unverified against the book (a library task).
- **Where, cost, effort.** $0; 12 person-days; 4–6 weeks.
- **Impact 3 (OB).** Risk: the firewall must be explained to a new audience; no value profile may be implied. **ROI 2.5.**

W15. Cross-cultural condemnation (fold into W9)

- **Objective.** Condemnation of hypocrisy is weaker in interdependent cultures (Effron et al. 2018). Quota-sample W9 across three cultural settings (N = 300 each) so that the attribution result is not a WEIRD-sample artefact.
- **Where, cost, effort.** Adds ≈ $1,300 and 3 person-days to W9. **Impact 3 (MP). ROI 1.0** on its own; run only as W9's design.

W16. A registered report in dialogue with Hua (2026) (BI §7G(3); W7 and W10)

- **Objective.** Hua's preregistration predicts no wiggle-room effect in models without a self-image; the lane's record shows compromise-seeking without one. Combine W7's hidden-consequences variant and W10's Dana replication into a registered report whose predictions are written before the arms run, so the result cannot be read as post hoc.
- **What exists.** The hypothesis and the counterpoint (B1 I-9); no preregistration.
- **Where, cost, effort.** $0 beyond W7 and W10; 5 person-days of writing; a registered-report venue's Stage 1 timeline (8–16 weeks).
- **Impact 4 (MP, ML).** Risk: Stage 1 review may require the expert gold arm first. **ROI 8.0.**

3. Summary table

Table 1. All workstreams. Cost is the expected metered or panel cost (cap in parentheses); effort in person-days; ROI by the §1 rule (impact × 10 ÷ (days + cash/$100)); venue codes MP, ML, OB as in §1. Order is by workstream id; the priority column gives the sequencing rank after dependencies.

id	workstream	corpus design	where	cost	effort (pd)	calendar	impact	venue	ROI	priority
W1	debrief audit: disclose / re-describe / misreport	§7A	owner machine	$0	3	1 wk	5	MP, ML	16.7	1
W11	reversed-priority twins	§S25; council 03	metered	$8 ($10)	3	1–2 wk	4	MP, ML	12.9	2
W12	structured intended-action field	§7K	metered	$8 ($10)	1.5	1 wk	3	ML	18.8	3
W8	post-action re-description; self vs other	§7I, §7J	metered	$16 ($20)	2	1–2 wk	4	MP	18.5	4
W3	salience arms: restatement vs comparison	§7C	metered	$16 ($20)	3	1–2 wk	4	MP	12.5	5
W13	expert gold arm resumed, with disclosure ratings	C1; A-26	humans	$1,000–2,000	5	3–4 wk	5	MP	2.5	6 (start now; long lead)
W6	analysis ablation × observation	§7F	metered	$32 ($40)	3	2 wk	4	MP	11.8	7
W7	menu manipulation (remove middle; deferral; hidden consequences)	§7G	metered	$24 ($30)	6	3 wk	5	MP, ML	6.8	8
W2	disclosure arm	§7B	metered	$8 ($10)	2	1 wk	3 (4 with W9)	MP	14.3	9 (after W3)
W4	pressure decomposition, three single-factor twins	§7D; §S25	metered	$24 ($30)	9	3–4 wk	4	MP, ML	3.9	10
W5	cost gradient, three steps	§7E	metered	$24 ($30)	8	3 wk	4	MP	4.3	11 (merge with W4)
W10	cross-species base-rate arm (coin, wiggle room, die roll)	new	metered	$10 ($15)	5	2 wk	4	MP, ML	6.6	12
W9	human observer study (+ W15 culture quota)	§7L	humans	$2,500 (+$1,300)	20 (+3)	8–12 wk	5	MP	1.1	13 (start ethics now)
W16	registered report with W7(3) and W10	§7G	writing	$0	5	8–16 wk	4	MP, ML	8.0	14
W14	organizational-decoupling paper	§7.3	writing	$0	12	4–6 wk	3	OB	2.5	15
W15	cross-cultural condemnation	§7L	humans	$1,300	3	with W9	3	MP	1.0	with W9

Totals. Metered arms W2–W8, W10–W12: about **$170 expected, $215 cap**, or roughly the four gap-closure arms again; human work W9, W13, W15: about **$5,000–6,000** including honoraria; effort about **90 person-days** across everything, of which the two authoring campaigns (W4, W5, W7, W10, W11: about 31 person-days) are the critical resource. The minimal set that decides whether the word may be used is **W1 + W11 + W8 + W3** (about $40, 11 person-days, three weeks); the minimal set that makes the result publishable to a moral-psychology audience adds **W13 and W9** (about $4,000–5,000 and three months of lead time).

4. Sequencing

The dependencies are few and sharp. W1 is $0 and gates the word; W11 can retire the framing and must run before any arm that presupposes the charter does work; W3 should precede W2 because a disclosure instruction is a comparison prompt; W9 needs W1's codes for its disclosed/undisclosed cells and W13's readers can rate disclosure if W1 has run; W16 needs W7(3) and W10 as its Stage 1 methods. Everything else is parallel and limited only by authoring time and the two-driver rate limit on the metered route (`GAP_CLOSURE_PLAN` §2).

Table 2. A twelve-week schedule with three metered grants and two human tracks. Weeks are calendar weeks from the first grant; "owner" means the owner's machine or decision.

week	$0 track (cloud and owner)	metered track (grants)	human track
1	W1 debrief audit (owner machine); W11 and W12 authored and validated; W8 and W3 instruction variants written	—	W13 readers approached; W9 ethics application drafted
2	W1 report; B2 wording decided by its result	grant 1 (≈ $50): W11, W12 + W8 shared administration, W3 two arms	W13 packet with W1 codes beside each record
3	W7 and W10 authoring begins	W11 result: continue or retire the framing; W3 result: mirror or labelled coin	W13 reads in progress
4	W7 variants validated; predicate audit extended	grant 2 (≈ $60): W6 four arms; W2 if W3 warrants	W13 ratings returned; A-26 STOP-and-ask
5–6	W4 (or the merged W4+W5 factorial) authoring; W10 validated	W7 three variants; W10	W9 preregistration written from W1 and W13
7–8	W16 Stage 1 drafted from W7(3) and W10 designs	grant 3 (≈ $60): W4 (or factorial)	W9 pilot (N = 100) after ethics approval
9–10	W14 outline with an organization theorist	W4 results; cost-curve figure	W9 main sample (N = 600, or 900 with W15)
11–12	v2 of the B2 paper assembled from W1, W3, W6, W7, W11	—	W9 analysis; W13 decision rule executed

If only some of it lands. W1 alone settles whether "hypocrisy" can appear; W1 + W11 settle whether there is a charter–action gap to name; W3 + W6 + W8 turn the Batson taxonomy from analogy into a set of tested mechanisms at under $70; W13 + W9 are what a psychology venue will ask for and cannot be bought faster than their lead times; W7 and W10 are the two arms that make the result a contribution to psychology rather than a borrowing from it.

References

Alicke, M., Gordon, E. & Rose, D. (2013). Hypocrisy: what counts? *Philos. Psychol.*, https://davdrose.github.io/assets/pdf/Hypocrisy-%20what%20counts.pdf · Barden, J., Rucker, D. D. & Petty, R. E. (2005). *PSPB*, https://doi.org/10.1177/0146167205276430 · Bardi, A. & Schwartz, S. H. (2003). Values and behavior. *PSPB*, https://journals.sagepub.com/doi/10.1177/0146167203254602 · Batson, C. D. et al. (1999). *JPSP*, https://pubmed.ncbi.nlm.nih.gov/10510506/ · Batson, C. D. & Thompson, E. R. (2001). *Curr. Dir. Psychol. Sci.*, https://api.crossref.org/works/10.1111/1467-8721.00114 · Dana, J., Weber, R. A. & Kuang, J. X. (2007). *Econ. Theory*, https://www.scheller.gatech.edu/directory/research/accounting/kuang/pdf/dwk_2007.pdf · Ditto, P. H., Pizarro, D. A. &

Tannenbaum, D. (2009). Motivated moral reasoning. *Psychol. Learn. Motiv.* · Dong, M. et al. (2019). *PLOS ONE*, https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0219382 · Effron, D. A. et al. (2018). Hypocrisy and culture. *JESP*, https://lbsresearch.london.edu/id/eprint/950 · Effron, D. A., O'Connor, K., Leroy, H. & Lucas, B. J. (2018). From inconsistency to hypocrisy. *Res. Organ. Behav.*, https://api.openalex.org/works/doi:10.1016/j.riob.2018.10.003 · Hua, T. (2026). AEA RCT Registry, https://www.socialscienceregistry.org/versions/303697/docs/version/document · Jordan, J. J. et al. (2017). *Psychol. Sci.*, https://www.psychologicalscience.org/journals/psychological-science/0956797616685771/ · Lynch, A. et al. (2025). Agentic misalignment, https://arxiv.org/abs/2510.05179 · Monin, B. & Miller, D. T. (2001). Moral credentials. *JPSP*, https://sites.nd.edu/lapseylab/files/2015/09/10.1.1.462.8290.pdf · O'Connor, K., Effron, D. A. & Lucas, B. J. (2020). *JPSP*, https://lbsresearch.london.edu/id/eprint/1398/ · Rotella, A. et al. (2026). https://researchportal.northumbria.ac.uk/en/publications/observation-moderates-the-moral-licensing-effect-a-meta-analytic-/ · Rustichini, A. & Villeval, M. C. (2014). *JEBO*, https://docs.iza.org/dp6590.pdf · Schoen, B. et al. (2025). Anti-scheming training, https://arxiv.org/abs/2509.15541 · Shalvi, S. et al. (2011). Justified ethicality. *OBHDP*, https://ideas.repec.org/a/eee/jobhdp/v115y2011i2p181-190.html · Valdesolo, P. & DeSteno, D. (2008). *JESP*, https://desteno.squarespace.com/s/the-duality-of-virtue-deconstructing-the-moral-hypocrite.pdf · Vu, L. et al. (2023). *Psychol. Bull.*, https://research.tilburguniversity.edu/en/publications/ignorance-by-choice-a-meta-analytic-review-of-the-underlying-moti/ · Wang (2026). Faithfulness of extracted conclusions, https://arxiv.org/abs/2606.00476. Lane: `plans/GAP_CLOSURE_PLAN_20261028.md`, `plans/SPEND_LEDGER.md`, `TASKS.md`, `PAPER_PLAN.md` §5, `paper/PAPER_v22.md` §§S25, `audits/EXPERT_READ_PROTOCOL_20261027.md`.